
Verify-or-Abstain Evidence Pipelines: When Cryptographic Custody Is Necessary but Not Sufficient

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Scientific evidence pipelines increasingly rely on content-addressed custody and
2 hash-valid artifacts, yet syntactically valid objects can remain semantically in-
3 consistent. We present an evidence-pipeline reliability study centered on explicit
4 terminal states, row-level accounting invariants, and closed-vocabulary identity
5 guards. Preregistered audits show that (i) a historical aggregate reported zero ab-
6 stentions while row-level data contained three; (ii) a successor implementation
7 with derived aggregates restores invariant consistency; (iii) a deterministic iden-
8 tity contract passes under synthetic fixtures but a single observed-source pair ad-
9 mits downstream passage when a guard is bypassed; and (iv) historical sequence
10 and copy-number corpora could not be recovered from admitted provenance. A
11 morphology null benchmark is retained as disclosed negative cached evidence.
12 We do not claim therapeutic efficacy, clinical utility, population-level biomedical
13 error rates, or real-world evidence.

14 1 Introduction

15 Machine-learning systems for biomedical hypothesis generation must separate *what was computed*
16 from *what may be claimed*. Content-addressed storage and hash verification establish integrity of
17 bytes, not meaning: aggregates can disagree with row ledgers, evaluation wiring can fail to exercise
18 intended disagreements, and missing corpora can leave historical numbers unverified. We frame this
19 work as an evidence-pipeline control problem: normalize, validate, admit or abstain, derive results,
20 and attach claim ceilings before projection to manuscripts [? ?].

21 2 Related Work and Terminology

22 Selective prediction and reject options allow models to abstain under uncertainty [? ?]. Workflow
23 and data provenance standards document lineage without guaranteeing semantic correctness [?].
24 Biomedical identifier normalization relies on authorities such as HGNC and UniProt [? ?]. Re-
25 producibility guidance distinguishes replication from reproducibility and encourages reporting null
26 results [?]. Regulatory Real-World Evidence frameworks define boundaries we do not claim here
27 [?].

28 We distinguish *cryptographic custody* (hash-valid, content-addressed storage) from *semantic verifi-*
29 *cation* (terminal-state invariants, identity resolution, aggregate-row consistency).

30 **Synthetic fixture.** An intentionally constructed input used to test deterministic software behavior.

31 **Canonical contract fixture.** A project-defined synthetic or curated fixture whose expected output
32 is fixed before execution and used to test a deterministic contract.




Figure 1: Verify-or-abstain evidence pipeline. Custody without semantic checks can admit inconsistent aggregates.

Experiment	Unit	Guard	Condition	N	Ceiling
GAP accounting repair	row	aggregate invariant	successor run	4	semantic audit
G3 contract	fixture	identity	Lane A	synth.	deterministic
G3 source ablation	pair	identity bypass	Lane B	1	descriptive
G1 provenance	corpus	—	not recovered	—	audit only
G2 provenance	corpus	—	not recovered	—	audit only
Morphology null	ranking	shuffle null	cached	50	repurposing hyp.

Table 1: Primary results from frozen experiment artifacts. Historical G1/G2 counts are excluded.

Observed source record. An actual record retrieved from an identified external source rather than invented for testing. Evaluations using such records are *observed-source evaluations*; they are not Real-World Evidence (RWE) unless regulatory RWD/RWE criteria are satisfied [?].

External biological evidence. A source-backed biological record or publication used as evidence for a biological assertion.

Abstention / reject option. Following Chow [?] and selective-prediction literature [?], a pipeline may emit a terminal ABSTAIN state rather than admit an under-supported record.

Closed-vocabulary identity resolution. Map source strings to a frozen alias table with exact, mapped, or unresolved outcomes; unresolved inputs must not silently pass as canonical.

Provenance. Origin and processing lineage of entities [?]. Hash equality does not imply semantic consistency between derived artifacts.

3 Verify-or-Abstain Architecture

Figure ?? sketches the pipeline: SOURCE → normalize/validate → admit or abstain → derived result → claim with ceiling → frozen custody graph. Semantic verification (terminal states, invariants, identity guards) complements cryptographic custody.

4 Experimental Design

Experiments are preregistered where noted; August 13 historical runs are preserved and not retroactively relabeled. EXP-GAP-ACCOUNTING-001.1 repairs aggregate derivation from rows. EXP-004 ablates a closed-vocabulary identity guard (Lane A: contract; Lane B: $N=1$ observed-source pair). EXP-002/003 provenance audits and external recovery (EXP-002-PROV.1, EXP-003-PROV.1) search admitted custody without rerunning frozen closeouts. EXP-005 preregistration remains locked but corpus/metric incomplete (not executed). EXP-006 documents a cached morphology null benchmark.

5 Results

Table ?? summarizes terminal classifications and claim ceilings from machine-readable receipts.

GAP accounting. Historical summary reported zero abstentions; row-level candidates contained three ABSTAIN terminals. Successor run: $N_{input}=4$, $N_{admitted}=1$, $N_{abstained}=3$, summary abstain count = 3; invariants pass.

Identity guard. Lane A: deterministic closed-vocabulary contract passes. Lane B: one observed source/canonical pair (TAZ/cardioliipin module → TAZ) yields ABSTAIN under guard

63 and downstream admission under bypass (*unsupported identity admission*); not a population error
64 rate. Historical evaluation invoked `reconcile(symbol, symbol)`; regression confirms the wiring
65 defect.

66 **Provenance recovery.** Admitted external bootstrap package (206 MB, SHA-256 verified) does not
67 contain recoverable G1/G2 derivation chains; numeric substring matches in unrelated CSV fields
68 are rejected as evidence.

69 **Null benchmark.** Cached evaluation: known-pair rank 28/50, reciprocal-rank enrichment vs. shuf-
70 fle < 1 ; disclosed miss not suppressed.

71 6 Discussion and Limitations

72 Scope is intentionally small: $N=1$ observed-source ablation; G1/G2 inputs not recovered; same
73 designers performed parts of evaluation; no therapeutic efficacy, biological rescue, clinical utility,
74 RWE, or population prevalence claims; EXP-005 replication not executed pending corpus freeze;
75 infrastructure integrity does not establish semantic correctness.

76 7 Reproducibility and Provenance

77 Artifacts, preregistrations, and SHA-256 manifests reside in the project Frozen Custody Graph.
78 Anonymous supplementary materials omit identifying metadata. Historical unverified counts remain
79 quarantined as UNVERIFIED_HISTORICAL_CLAIM.

80 8 Conclusion

81 Hash-valid custody is necessary but insufficient for trustworthy evidence pipelines. Explicit termi-
82 nal states and invariants surface failure modes that remain syntactically admissible—a prerequisite
83 before any biomedical performance claim.

84 References

85 NeurIPS Paper Checklist

86 The checklist is designed to encourage best practices for responsible machine learning research,
87 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not re-
88 move the checklist: **The papers not including the checklist will be desk rejected.** The checklist
89 should follow the references and follow the (optional) supplemental material. The checklist does
90 NOT count towards the page limit.

91 Please read the checklist guidelines carefully for information on how to answer these questions. For
92 each question in the checklist:

- 93 • You should answer blue[Yes], orange[No], or gray[N/A].
- 94 • gray[N/A] means either that the question is Not Applicable for that particular paper or the
95 relevant information is Not Available.
- 96 • Please provide a short (1–2 sentence) justification right after your answer (even for
97 gray[N/A]).

98 **The checklist answers are an integral part of your paper submission.** They are visible to the
99 reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it
100 (after eventual revisions) with the final version of your paper, and its final version will be published
101 with the paper.

102 The reviewers of your paper will be asked to use the checklist as one of the factors in their eval-
103 uation. While blue[Yes] is generally preferable to orange[No], it is perfectly acceptable to answer
104 orange[No] provided a proper justification is given (e.g., error bars are not reported because it would
105 be too computationally expensive” or “we were unable to find the license for the dataset we used”).

106 In general, answering orange[No] or gray[N/A] is not grounds for rejection. While the questions are
107 phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just
108 use your best judgment and write a justification to elaborate. All supporting evidence can appear
109 either in the main paper or the supplemental material, provided in appendix. If you answer blue[Yes]
110 to a question, in the justification please point to the section(s) where related material for the question
111 can be found.

112 IMPORTANT, please:

- 113 • **Delete this instruction block, but keep the section heading “NeurIPS Paper Check-**
114 **list”,**
- 115 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 116 • **Do not modify the questions and only use the provided macros for your answers.**

117 1. Claims

118 Question: Do the main claims made in the abstract and introduction accurately reflect the
119 paper’s contributions and scope?

120 Answer: red[TODO]

121 Justification: red[TODO]

122 Guidelines:

- 123 • The answer gray[N/A] means that the abstract and introduction do not include the
124 claims made in the paper.
- 125 • The abstract and/or introduction should clearly state the claims made, including the
126 contributions made in the paper and important assumptions and limitations. A or-
127 ange[No] or gray[N/A] answer to this question will not be perceived well by the re-
128 viewers.
- 129 • The claims made should match theoretical and experimental results, and reflect how
130 much the results can be expected to generalize to other settings.
- 131 • It is fine to include aspirational goals as motivation as long as it is clear that these
132 goals are not attained by the paper.

133 2. Limitations

134 Question: Does the paper discuss the limitations of the work performed by the authors?

135 Answer: red[TODO]

136 Justification: red[TODO]

137 Guidelines:

- 138 • The answer gray[N/A] means that the paper has no limitation while the answer or-
139 ange[No] means that the paper has limitations, but those are not discussed in the paper.
- 140 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 141 • The paper should point out any strong assumptions and how robust the results are to
142 violations of these assumptions (e.g., independence assumptions, noiseless settings,
143 model well-specification, asymptotic approximations only holding locally). The au-
144 thors should reflect on how these assumptions might be violated in practice and what
145 the implications would be.
- 146 • The authors should reflect on the scope of the claims made, e.g., if the approach was
147 only tested on a few datasets or with a few runs. In general, empirical results often
148 depend on implicit assumptions, which should be articulated.
- 149 • The authors should reflect on the factors that influence the performance of the ap-
150 proach. For example, a facial recognition algorithm may perform poorly when image
151 resolution is low or images are taken in low lighting. Or a speech-to-text system might
152 not be used reliably to provide closed captions for online lectures because it fails to
153 handle technical jargon.
- 154 • The authors should discuss the computational efficiency of the proposed algorithms
155 and how they scale with dataset size.

- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper does not include experiments.
- If the paper includes experiments, an orange[No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.

- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so orange[No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: red[TODO]

Justification: red[**TODO**]

Guidelines:

- The answer gray[N/A] means that the paper does not include experiments.
- The authors should answer blue[Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: red[**TODO**]

Justification: red[**TODO**]

Guidelines:

- The answer gray[N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: red[**TODO**]

Justification: red[**TODO**]

Guidelines:

- The answer gray[N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer orange[No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that there is no societal impact of the work performed.
- If the authors answer gray[N/A] or orange[No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.

- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: red[TODO]

Justification: red[TODO]

Guidelines:

- The answer gray[N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- 416 • Depending on the country in which research is conducted, IRB approval (or equivalent)
417 may be required for any human subjects research. If you obtained IRB approval,
418 you should clearly state this in the paper.
- 419 • We recognize that the procedures for this may vary significantly between institutions
420 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
421 guidelines for their institution.
- 422 • For initial submissions, do not include any information that would break anonymity
423 (if applicable), such as the institution conducting the review.

424 16. Declaration of LLM usage

425 Question: Does the paper describe the usage of LLMs if it is an important, original, or
426 non-standard component of the core methods in this research? Note that if the LLM is used
427 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
428 scientific rigor, or originality of the research, declaration is not required.

429 Answer: red[TODO]

430 Justification: red[TODO]

431 Guidelines:

- 432 • The answer gray[N/A] means that the core method development in this research does
433 not involve LLMs as any important, original, or non-standard components.
- 434 • Please refer to our LLM policy in the NeurIPS handbook for what should or should
435 not be described.